

Exploring a million songs

Parallel Regression Trees in Ufora



UFORA

Ronen Hilewicz, VP of Engineering
Ronen@Ufora.com

Agenda

- Ufora – language and distributed runtime
- Million Song Database
- The Problem - Year Prediction
- Prior Approaches
- Regression Trees Refresher
- Demo
- Performance and Scalability

What is **U**FORA?

An *implicitly* parallel language and runtime for data science

IMPLICIT PARALLELISM

Code that would normally
execute serially...



... in Ufora will **split** across multiple
processes on multiple machines, ...



... use more data, and finish faster.

Million Song Database

- A 280GB free collection of audio features and metadata for a million contemporary popular music tracks
- Does not contain any audio – only the derived features
- Provided by Columbia's LabROSA
(Recognition and Organization of Speech and Audio)
- Created by The Echo Nest



The Data

Static Metadata:

Title • Artist • Album • Genre • Year • Artist Location • etc.

Dynamic Metadata (from EchoNest API):

Song/Artist Hottness • Familiarity • Similar Artists • Artist Keywords

Audio Features:

- Key • Mode • Time-Signature • Tempo
- Loudness • Energy
- sections [bars [beats [tatums [segments


Demo

Exploring the Dataset

Year Prediction

- From audio features, try to predict the year when the song was released

- Timbre:**



Year Prediction - Prior Approaches

- **The Million Song Dataset**, T. Bertin-Mahieux, D. Ellis, B. Whitman and P. Lamere, ISMIR '11
- Methods:
 - Baseline: Uniform prediction
 - k nearest neighbors (k -NN)
 - Vowpal Wabbit – linear regression using gradient descent
- Challenges:
 - Only half the dataset has year info


Prior Results

Constant	8.13	10.80
1-NN	9.81	13.99
50-NN	7.58	10.20
VW	6.14	8.76

Our Approach

- Regression Trees
 - Highly parallelizable
 - Have been implemented at scale using MapReduce (PLANET)
 - Recently implemented a regression and classification tree library in Ufora

Regression Trees Refresher



Learning The Model

- Using a greedy top-down approach
- Partition D^* along the split predicate, and proceed recursively on the partitions to build child nodes.
- We select split predicates that minimize the *impurity* in Y values of the training records that are passed to the node. In other words, we want to maximize

$$|D| \times \text{Var}(YD) - (|DL| \times \text{Var}(YDL) + |DR| \times \text{Var}(YDR))$$

Where YDL and YDR are the Y values in the partitions DL and DR of D


Demo

Building the Model

Results


Constant	8.13	10.80
1-NN	9.81	13.99
50-NN	7.58	10.20
Tree Model	6.87	9.66
VW	6.14	8.76

Performance (in seconds)



Depth	SciKit	Ufora
1	7.02	1.4
5	35.14	6.43
8	56.21	10.7
10	85.16	17.18

Comparison of Approaches



Resources

<http://labrosa.ee.columbia.edu/millionsong/>

Thierry Bertin-Mahieux, Daniel P.W. Ellis, Brian Whitman, and Paul Lamere.
The Million Song Dataset. In Proceedings of the 12th International Society
for Music Information Retrieval Conference (ISMIR 2011), 2011.

- Vowpal Wabbit - https://github.com/JohnLangford/vowpal_wabbit/wiki
- SciKit Learn - <http://scikit-learn.org/>

Biswanath Panda, Joshua S. Herbach, Sugato Basu, Roberto J. Bayardo

PLANET: Massively Parallel Learning of Tree Ensembles with MapReduce

Proceedings of the 35th International Conference on Very Large Data Bases (VLDB)

We're Hiring!

- Lead Front End Engineer – Javascript, node.js, Angular.js
- Senior C++ Engineer – C++, boost, Python

info@ufora.com

ronen@ufora.com